Synthetic DNA Motif Classifier

Biology / Sequence ML | Implementation and interview guide

Purpose and intended demonstration

Learn short-sequence statistics for an artificial motif recognition task without functional genomic claims. This is a computational biology learning project with synthetic or fictional inputs, not a clinical or experimentally validated system.

How the implementation works

Generate random DNA strings and insert a short arbitrary tag in positive examples. Encode normalized dinucleotide frequencies and train a softmax classifier. Compare with direct tag matching to expose when a simple rule is better.

Evaluation and evidence

Use independent generated sequence sets for training and test. Report ML classification metrics and exact-tag baseline performance.

Synthetic DNA Motif Classifier

Biology / Sequence ML | Implementation and interview guide

Setup and commands

Requires Python 3.10 or newer and NumPy. Create a virtual environment: python -m venv .venv. On Windows use .venv\Scripts\python.exe if activation is unavailable. Install with python -m pip install -r requirements.txt. Run python app.py, then open results/report.html in your browser. Run python -m unittest discover -v to execute tests. The default seed is 42; change it with --seed. Replace sample_input.json using --input path-to-file.json, preserving its schema. Algorithm-only projects generate synthetic arrays from the seed rather than consuming external files.

Architecture and file responsibilities

CLI arguments -> JSON input or seeded generator -> domain pipeline -> validation or evaluation -> JSON and HTML artifacts. Imported functions can be used in a notebook or a separate application. Model checkpoints preserve learned arrays where applicable; see app.py for parameter order. This package does not include a production web service. Each project is independent and includes its own helper code so the ZIP can be extracted and run alone. app.py holds the experiment and domain functions: validate_sequence, dinucleotides, generate, run. kit.py contains reporting, input loading and only the retrieval or generation helpers needed by this project. sample_input.json contains editable inputs. test_app.py checks domain behaviors; test_reporting.py checks metrics and HTML output boundaries. results/metrics.json records outputs and results/report.html provides a browser-readable report. Numeric projects also save arrays or learned parameters to results/artifacts.npz; inspect its keys with numpy.load. MODEL_CARD.md explains intended use and limits. .github/workflows/tests.yml runs tests and the fixture experiment. No API key is needed.

Synthetic DNA Motif Classifier

Biology / Sequence ML | Implementation and interview guide

Limitations and next improvements

The tag is an artificial label, not a promoter, binding site or pathogenic motif. Dinucleotide counts lose positional information and may underperform direct matching. Use a licensed regulatory-sequence benchmark, chromosome or homology-aware splits, reverse-complement handling and position-aware CNN or attention baselines.

Interview walkthrough

Present the problem and demonstrate one normal example, then one failure or boundary case. Explain the baseline or fixture mode, the split or validation design, and one limitation before discussing an extension. Be ready to trace data through the functions listed above. Explain why the chosen metric matches the task and why the available results cannot establish production performance. Add your own documented experiment and be precise about what you implemented or changed. These are AI-assisted educational projects; filenames and synthetic scores are not evidence of client work.

Research context and reading

These topics were selected from research directions reviewed on 22 October 2026. Recent activity suggests useful skills to develop; it is not a forecast of hiring, commercial success or scientific validity. The implementation is a small educational baseline and does not reproduce or claim the performance of any cited foundation model. Reading: AlphaGenome regulatory prediction research (2026): https://www.nature.com/articles/s41586-025-10014-0.

Verified build results

project

"Synthetic DNA Motif Classifier"

ml metrics

{"accuracy": 0.7366666666666667, "precision": 0.7697368421052632, "recall": 0.7267080745341615, "f1": 0.7476038338658146, "tp": 117, "fp": 35, "fn": 44, "tn": 104}

exact tag baseline

{"accuracy": 0.99, "precision": 0.9817073170731707, "recall": 1.0, "f1": 0.9907692307692307, "tp": 161, "fp": 3, "fn": 0, "tn": 136}

tag

"ACGTAC"

scope

"Artificial sequence classification; no functional genomic inference"

See results/metrics.json for all observations and TEST_RESULTS.txt for the executed tests. All datasets and replay responses shipped here are synthetic. Results can vary across seeds and numerical library versions.